

μ PandOS: Phase 1

Luca Bassi (luca.bassi14@studio.unibo.it)

Gabriele Genovese (gabriele.genovese2@studio.unibo.it)

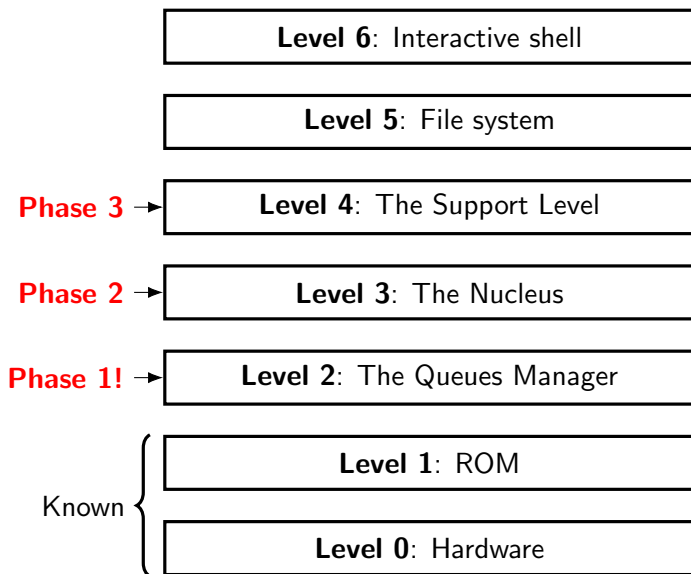
November 6, 2023

- ▶ μ PandOS is an educational *microkernel* operating system
- ▶ evolution of AMIKaya and PandOS; they are also the evolution of a long list of O.S. proposed for educational purposes (PandOS+, Kaya, HOCA, TINA, ICARO, etc).
- ▶ μ PandOS may be realized on the μ MPS3 architecture or on the μ RISCV architecture (for the brave's ones)
- ▶ Architecture based on six levels of abstraction, on the model of the O.S. THE proposed by Dijkstra in one of his articles of 1968

- ▶ Emulates a MIPS architecture (introduced in 1985)
- ▶ MIPS architecture is widely used in embedded systems
- ▶ Tool's version 3 released in 2020
- ▶ Pros: tool tested and working, easy to install
- ▶ Cons: limited debugger

- ▶ Emulates a RISC-V architecture (project started in 2010 and it's evolving)
- ▶ RISC-V architecture is fully open source
- ▶ Tool not yet fully tested and released
- ▶ Pros: debug with GDB
- ▶ Cons: no GUI (WIP), tool and toolchain need to be compiled from source

μ PandOS: six levels of abstraction



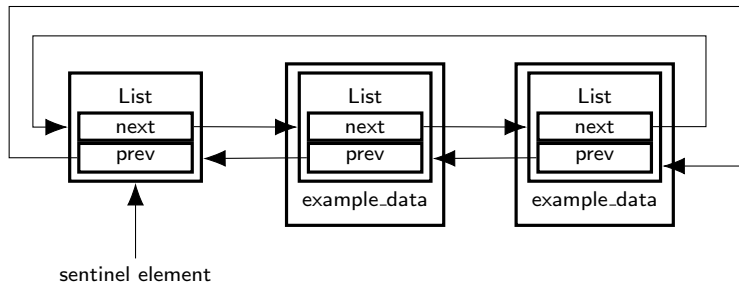
μ PandOS Phase 1 requests

- ▶ For this first phase, you have to add the functions requested, and the test procedure must conclude correctly.
- ▶ Level 2 requirements are purely logical; it is not yet touched the specific hardware of the emulator.
- ▶ You are provided with a **.tar.gz** file, that contains useful files to begin with.

The listx module

- ▶ A provided module, with a subset of functions of the Linux's `list.h` module
- ▶ It implements generic and type-oblivious lists
- ▶ Take a closer look to its content!!

```
struct list_head {  
    struct list_head *next;  
    struct list_head *prev;  
}
```



Level 2 specs: PCB

- ▶ μ PandOS level 2 (The Queues Manager) provides the implementation of the data structures used from the level 3 (The Nucleus)
- ▶ The Process Control Block (PCB) represent a process in the system.
- ▶ The queue manager implements PCB-related features:
 - Allocation and Deallocation of PCBs.
 - PCB Queue Management.
 - PCB Tree Management.
 - Management of incoming messages.
- ▶ ASSUMPTION: no more than MAXPROC (40) concurrent processes in μ PandOS

PCB data structure

```
/* process table entry type */
typedef struct pcb_t
{
    /* process queue */
    struct list_head p_list;

    /* process tree fields */
    struct pcb_t *p_parent; /* ptr to parent */
    struct list_head p_child; /* children list */
    struct list_head p_sib; /* sibling list */

    /* process status information */
    state_t p_s; /* processor state */
    cpu_t p_time; /* cpu time used by proc */

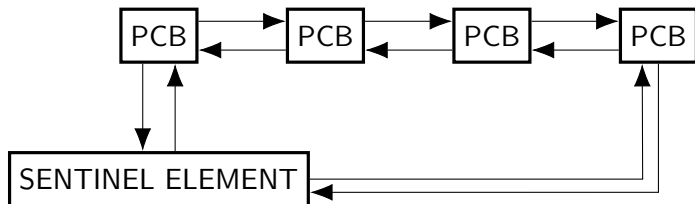
    /* First message in the message queue */
    struct list_head msg_inbox;

    /* Pointer to the support struct */
    support_t *p_supportStruct;

    /* process id */
    int p_pid;
} pcb_t, *pcb_PTR;
```

PCB's list

- ▶ PCBs can be organized into queues, called process queues (e.g. active process queue).
- ▶ Each list is managed via list_head fields.
- ▶ A list is identified by a sentinel element of type list_head.



PCB allocation

Main data structures:

- ▶ `pcbFree_h`: free or not used PCBs' list.
- ▶ `pcbTable[MAXPROC]`: array of PCBs with max size of MAXPROC.

PCB's functions to implement

Allocation and Deallocation of PCBs:

1. `void freePcb(pcb_t *p)`: insert the element pointed to by `p` onto the `pcbFree` list.
2. `pcb_t *allocPcb()`: return NULL if the `pcbFree` list is empty. Otherwise, remove an element from the `pcbFree` list, provide initial values for ALL of the PCBs fields and then return a pointer to the removed element. PCBs get reused, so it is important that no previous value persist in a PCB when it gets reallocated.
3. `void initPcbs()`: initialize the `pcbFree` list to contain all the elements of the static array of `MAXPROC` PCBs. This method will be called only once during data structure initialization.

PCB's functions to implement

PCB Queue Management:

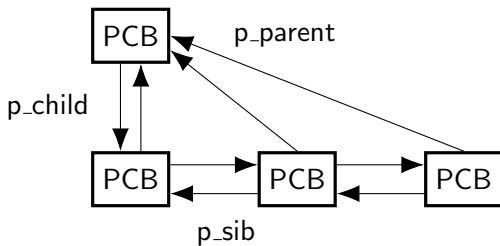
4. `void mkEmptyProcQ(struct list_head *head)`: this method is used to initialize a variable to be head pointer to a process queue.
5. `int emptyProcQ(struct list_head *head)`: return TRUE if the queue whose head is pointed to by head is empty. Return FALSE otherwise.
6. `void insertProcQ(struct list_head *head, pcb_t *p)`: insert the PCB pointed by p into the process queue whose head pointer is pointed to by head.

PCB's functions to implement

7. `pcb_t *headProcQ(struct list_head *head)`: return a pointer to the first PCB from the process queue whose head is pointed to by head. Do not remove this PCB from the process queue. Return NULL if the process queue is empty.
8. `pcb_t *removeProcQ(struct list_head *head)`: remove the first (i.e. head) element from the process queue whose head pointer is pointed to by head. Return NULL if the process queue was initially empty; otherwise return the pointer to the removed element.
9. `pcb_t *outProcQ(struct list_head *head, pcb_t *p)`: remove the PCB pointed to by p from the process queue whose head pointer is pointed to by head. If the desired entry is not in the indicated queue (an error condition), return NULL; otherwise, return p. Note that p can point to any element of the process queue.

PCB's trees

- ▶ In addition to the possibility of participating in a process queue, PCBs can be organized into process trees.
- ▶ Each parent contains a pointer to the list of the children (`p_child`)
- ▶ Each child has a pointer to the parent (`p_parent`) and a pointer that connects the brothers to each other.



PCB's functions to implement

PCB Tree Management:

10. `int emptyChild(pcb_t *p)`: return TRUE if the PCB pointed to by p has no children. Return FALSE otherwise.
11. `void insertChild(pcb_t *prnt, pcb_t *p)`: make the PCB pointed to by p a child of the PCB pointed to by prnt.
12. `pcb_t *removeChild(pcb_t *p)`: make the first child of the PCB pointed to by p no longer a child of p. Return NULL if initially there were no children of p. Otherwise, return a pointer to this removed first child PCB.
13. `pcb_t *outChild(pcb_t *p)`: make the PCB pointed to by p no longer the child of its parent. If the PCB pointed to by p has no parent, return NULL; otherwise, return p. Note that the element pointed to by p could be in an arbitrary position (i.e. not be the first child of its parent).

Messages

- ▶ Access to resources occurs through message passing
- ▶ The Message Block (**msg**) represent a message in the system.
- ▶ The queue manager implements msg-related features:
 - Allocation and Deallocation of msgs.
 - Msg Queue Management.
- ▶ ASSUMPTION: no more than MAXMESSAGES (40) concurrent messages in μ PandOS.

```
typedef struct msg_t
{
    /* message queue */
    struct list_head m_list;

    /* thread that sent this message */
    struct pcb_t *m_sender;

    /* the payload of the message */
    unsigned int m_payload;
} msg_t, *msg_PTR;
```

Messages allocation

Main data structures:

- ▶ `msgFree_h`: free or not used msgs' list.
- ▶ `msgTable[MAXMESSAGES]`: array of msgs with max size of MAXMESSAGES.

Msg's functions to implement

Allocation and Deallocation of msgs:

1. `void freeMsg(msg_t *m)`: insert the element pointed to by `m` onto the `msgFree` list.
2. `msg_t *allocMsg()`: return NULL if the `msgFree` list is empty. Otherwise, remove an element from the `msgFree` list, provide initial values for ALL of the messages fields and then return a pointer to the removed element. Messages get reused, so it is important that no previous value persist in a message when it gets reallocated.
3. `void initMsgs()`: initialize the `msgFree` list to contain all the elements of the static array of `MAXMESSAGES` messages. This method will be called only once during data structure initialization.

Msg's functions to implement

Msg Queue Management:

4. `void mkEmptyProcQ(struct list_head *head)`: used to initialize a variable to be head pointer to a message queue; returns a pointer to the head of an empty message queue.
5. `int emptyMessageQ(struct list_head *head)`: returns TRUE if the queue whose tail is pointed to by head is empty, FALSE otherwise.
6. `void insertMessage(struct list_head *head, msg_t *m)`: insert the message pointed to by m at the end (tail) of the queue whose head pointer is pointed to by head.

Msg's functions to implement

7. `void pushMessage(struct list_head *head, msg_t *m)`: insert the message pointed to by `m` at the head of the queue whose head pointer is pointed to by `head`.
8. `msg_t *popMessage(struct list_head *head, pcb_t *p_ptr)`: remove the first element (starting by the head) from the message queue accessed via `head` whose sender is `p_ptr`. If `p_ptr` is `NULL`, return the first message in the queue. Return `NULL` if the message queue was empty or if no message from `p_ptr` was found; otherwise return the pointer to the removed message.
9. `msg_t *headMessage(struct list_head *head)`: return a pointer to the first message from the queue whose head is pointed to by `head`. Do not remove the message from the queue. Return `NULL` if the queue is empty.

Suggestions

- ▶ There isn't a single way to implement Phase 1's functions.
- ▶ Use `static` methods and variables where possible.
- ▶ Learn to use a debugger: if you're using `uriscv`, use `gdb`; if you're using `umps3`, use the embedded one (it has some limitations).
- ▶ Get familiar with `umps3`

umps3 GUI

The screenshot displays the umps3 GUI, a simulator interface for a system. The main window is divided into several panes:

- Top Menu:** Simulator, View, Machine, Debug, Settings, Windows, Help.
- Overview Panel:** Shows system status (Stopped: User), suspect (50), type (Write), ASID (0x40), location (msgFree_h-0x...), and victims.
- Traced Regions:** Displays a memory dump for the region 0x2000.40ec, showing hex dump (Big-endian) and values (00 00 00 00).
- Assembly Code:** A list of instructions with addresses, opcodes, and operands. The code is for a process named "TermPrint-0x00" (Stopped: User). The instructions include:
 - 0x2000.1254: nop
 - 0x2000.1258: sw \$v0, +16(\$s8)
 - 0x2000.125c: beq \$s0, \$s0, 20001274
 - 0x2000.1260: nop
 - 0x2000.1264: lw \$s0, +24(\$s8)
 - 0x2000.1268: jal 2000114c <termstat+0x0>
 - 0x2000.126c: nop
 - 0x2000.1270: sw \$v0, +16(\$s8)
 - 0x2000.1274: lw \$v1, +16(\$s8)
 - 0x2000.1278: addiu \$v0, \$s0, +3
 - 0x2000.127c: beq \$v1, \$v0, 20001264
 - 0x2000.1280: nop
 - 0x2000.1284: lw \$v1, +16(\$s8)
 - 0x2000.1288: addiu \$v0, \$s0, +5
 - 0x2000.128c: beq \$v1, \$v0, 200012a4
 - 0x2000.1290: nop
 - 0x2000.1294: addiu \$v0, \$s0, +1
 - 0x2000.1298: sw \$v0, +28(\$s8)
 - 0x2000.129c: beq \$s0, \$s0, 200012b4
 - 0x2000.12a0: nop
 - 0x2000.12a4: lw \$v0, +48(\$s8)
 - 0x2000.12a8: nop
 - 0x2000.12ac: addiu \$v0, \$v0, +1
 - 0x2000.12b0: sw \$v0, +48(\$s8)
 - 0x2000.12b4: lw \$v0, +48(\$s8)
 - 0x2000.12b8: nop
 - 0x2000.12bc: lb \$v0, +0(\$v0)
 - 0x2000.12c0: nop
 - 0x2000.12c4: beq \$v0, \$s0, 20001304
 - 0x2000.12c8: nop
 - 0x2000.12cc: lw \$v0, +28(\$s8)
 - 0x2000.12d0: nop
 - 0x2000.12d4: beq \$v0, \$s0, 20001220
 - 0x2000.12d8: nop
 - 0x2000.12dc: beq \$s0, \$s0, 20001304
 - 0x2000.12e0: nop
 - 0x2000.12e4: addiu \$v0, \$s0, +1
 - 0x2000.12e8: nop
- Registers:** A table showing the values of CPU registers. The registers are:
 - Index: 0x00000000
 - Random: 0x00000000
 - EntryLo: 0x00000000
 - BadVAddr: 0x00000000
 - Timer: 0x00000000
 - EntryHi: 0x00000000
 - Status: 0x10000000
- Bottom Panel:** Shows the status (Stopped) and ToD (00000000.00000000.798).

Figure: Live Demo

Alternative tricks for debugging in umps3

- ▶ The provided file `klog.c` allows to “print” text and vars. On your code use the `klog_print` function, then on `umps3` add the `klog_buffer` as a traced region.
- ▶ Create an empty function to simulate a break-point

```
#include "klog.c"

void bp() {}

void scheduler() {
    ...
    klog_print("Sono arrivato qua\n");
    bp();
    ...
}
```


Setup tools

You need to install μ MPS3.

For example in Debian 12: `sudo apt install umps3`

This will also install the cross-compiler toolchain package:
`gcc-mipsel-linux-gnu`

Provided files

These file are provided in a **.tar.gz**:

- ▶ `const.h`: constant (`MAXPROC`, **`MAXMESSAGES`**, etc...) definition file
- ▶ `type.h`: struct types (`pcb_t`, `msg_t`, etc...) definition file
- ▶ `listx.h`: functions and macros for list manipulation
- ▶ `pcb.h`: PCB header file
- ▶ `msg.h`: message header file
- ▶ `pcb.c`: **write your implementation here**
- ▶ `msg.c`: **write your implementation here**
- ▶ `Makefile`: file to compile your project
- ▶ `p1test.c`: test file

Submission

The deadline is set for **Sunday December 31, 2023 at 23:59**.

Upload a single `phase1.tar.gz` in the folder associated to your group with:

- ▶ All the source code with a Makefile
- ▶ Documentation
- ▶ README and AUTHOR files

Please comment your code!

We will send you an email with the hash of the archive, you should check that everything is correct.

You will receive another email with the score out of 10.